

Quantization steps

trace code順序：

| | | |
|----|--------------------------|---|
| 1 | ristretto.cpp | quantize() |
| 2 | quantization.cpp | QuantizeNet() RunForwardBatches() |
| 3 | net.cpp | RangeInLayers() |
| 4 | quantization.cpp | Quantize2DynamicFixedPoint() EditNetDescriptionDynamicFixedPoint() Quantize2MiniFloat() EditNetDescriptionMiniFloat() Quantize2IntegerPowerOf2Weights() EditNetDescriptionIntegerPowerOf2Weights() |
| 5 | solver.cpp | Step() |
| 6 | net.hpp | ForwardBackward() |
| 7 | net.cpp | Forward(Dtype* loss) ForwardFromTo(int start, int end) |
| 8 | layer.hpp | Forward() |
| 9 | conv_ristretto_layer.cpp | Forward_cpu() |
| 10 | base_ristretto_layer.cpp | QuantizeLayerInputs_cpu() QuantizeWeights_cpu() QuantizeLayerOutputs_cpu() Trim2FixedPoint_cpu() Trim2MiniFloat_cpu() Trim2IntegerPowerOf2_cpu() |
| 11 | sgd_solver.cpp | ApplyUpdate() |

Dynamic Fixed Point

1. 從command line得到各種資料，用來initialize Quantization
2. 複製model和weight, net_val, 使用test data測試得到baseline accuracy

3. 複製model和weight, net_test, 使用train data來找出con.,ip. layer的
index(layer_names), 並算出該layer的in、out、param的最大值(max_in、max_out、
max_params)
此處param是weights
4. 對max_in(_out、_params)取log2, 可以得到integer length(所需bit數)
5. 從model讀出parameter, 並設定成test data
6. 對所有的con. layer的weight找出fl, 從16 bit width開始直到accuracy降到error margin
之外
 - a. 每跑完一種bit width, 會把accuracy儲存起來
 - b. 每個con. layer會用同一個bit width做測試
7. 對所有的ip. layer的weight找出fl, 從16 bit width開始直到accuracy降到error margin之
外
 - a. 每個ip. layer會用同一個bit width做測試
8. 對所有的con.、ip. layer的in、out找出fl, 從16 bit width開始直到accuracy降到error
margin之外
 - a. 每個layer的activation都用同一個bit width做測試
 - b. 6~8.不會同時Quantize, 一次只會對某個layer的某個part做Quantize後算出
accuracy
9. 將在error margin以內的accuracy, 選出最小bit width
 - a. bw_con、bw_fc、bw_out(bw_in = bw_out)
 - b. 每個con. layer都是用bw_con、ip. layer用bw_fc、兩者都使用bw_in、bw_out
10. 從model讀出parameter, 並設定成test data
11. 把bw_con、bw_fc、bw_out、bw_in放入EditNetDescriptionDynamicFixedPoint(), 得
到每個layer part都做過Quantize的net
12. 在開始training之後, Forwarding的過程中, 再做運算之前, 會先把input、output、
weight做quantize(Trim2FixedPoint_cpu)才開始該次forward

```

template <typename Dtype>
void BaseRistrettoLayer<Dtype>::Trim2FixedPoint_cpu(Dtype* data, const int cnt,
const int bit_width, const int rounding, int f1) {
    for (int index = 0; index < cnt; ++index) {
        // Saturate data
        Dtype max_data = (pow(2, bit_width - 1) - 1) * pow(2, -f1);
        Dtype min_data = -pow(2, bit_width - 1) * pow(2, -f1);
        data[index] = std::max(std::min(data[index], max_data), min_data);
        // Round data
        data[index] /= pow(2, -f1);
        switch (rounding) {
            case QuantizationParameter_Rounding_NEAREST:
                data[index] = round(data[index]);
                break;
            case QuantizationParameter_Rounding_STOCHASTIC:
                data[index] = floor(data[index] + RandUniform_cpu());
                break;
            default:
                break;
        }
        data[index] *= pow(2, -f1);
    }
}

```

- a. type仍然是32bit, 但Data的value已經被改為bw_con、bw_fc、bw_out、bw_in
可以表示的範圍

MiniFloat

1. 接續DFP步驟3.
2. 對max_in、max_out做兩次取log2，第一次是得到需要n次方，第二次是得到表達n次方所需的bit width，得exp_in、exp_out。接著取兩者間最大的作為exp_bit。
3. 換下一個layer做，得到新的exp_in、exp_out，若有更大的值就更新exp_bit
 - a. 因為weight通常小於activation，所以不另外計算
 - b. 所有layer都用一樣長度的exp_bit(?)
4. 從model讀出parameter，並設定成test data
5. 對所有layer的activation，從16 bit width開始測試直到accuracy降到error margin之外，每一次bitwidth的結果都會被記錄下來
 - a. input、output、weight都用同一個bit width
6. 在error margin之內選出最小的bit width，作為quantize net的結果

```

template <typename Dtype>
void BaseRistrettoLayer<Dtype>::Trim2MiniFloat_cpu(Dtype* data, const int cnt,
    const int bw_mant, const int bw_exp, const int rounding) {
    for (int index = 0; index < cnt; ++index) {
        int bias_out = pow(2, bw_exp - 1) - 1;
        float_cast d2;
        // This casts the input to single precision
        d2.d = (float)data[index];
        int exponent=d2.parts.exponent - 127 + bias_out;
        double mantisa = d2.parts.mantisa;
        // Special case: input is zero or denormalized number
        if (d2.parts.exponent == 0) {
            data[index] = 0;
            return;
        }
        // Special case: denormalized number as output
        if (exponent < 0) {
            data[index] = 0;
            return;
        }
    }
}

```

```

// Special case: denormalized number as output
if (exponent < 0) {
    data[index] = 0;
    return;
}
// Saturation: input float is larger than maximum output float
int max_exp = pow(2, bw_exp) - 1;
int max_mant = pow(2, bw_mant) - 1;
if (exponent > max_exp) {
    exponent = max_exp;
    mantisa = max_mant;
} else {
    // Convert mantissa from long format to short one. Cut off LSBs.
    double tmp = mantisa / pow(2, 23 - bw_mant);
    switch (rounding) {
        case QuantizationParameter_Rounding_NEAREST:
            mantisa = round(tmp);
            break;
        case QuantizationParameter_Rounding_STOCHASTIC:
            mantisa = floor(tmp + RandUniform_cpu());
            break;
        default:
            break;
    }
}
// Assemble result
data[index] = pow(-1, d2.parts.sign) * ((mantisa + pow(2, bw_mant)) /
    pow(2, bw_mant)) * pow(2, exponent - bias_out);
}
}

```

INTEGER_POWER_OF_2_WEIGHTS

1. weights直接指定exp_min = -8、exp_max = 1
2. Activations使用dynamix fix point, 用bitwidth=8

```
template <typename Dtype>
void BaseRistrettoLayer<Dtype>::Trim2IntegerPowerOf2_cpu(Dtype* data,
    const int cnt, const int min_exp, const int max_exp, const int rounding) {
    for (int index = 0; index < cnt; ++index) {
        float exponent = log2f((float)fabs(data[index]));
        int sign = data[index] >= 0 ? 1 : -1;
        switch (rounding) {
            case QuantizationParameter_Rounding_NEAREST:
                exponent = round(exponent);
                break;
            case QuantizationParameter_Rounding_STOCHASTIC:
                exponent = floorf(exponent + RandUniform_cpu());
                break;
            default:
                break;
        }
        exponent = std::max(std::min(exponent, (float)max_exp), (float)min_exp);
        data[index] = sign * pow(2, exponent);
    }
}
```